

An Introduction to Metamodelling and Graph Transformations

with eMoflon



Part II: Ecore

For eMoflon Version 2.15.0

File built on 3rd May, 2016

Copyright © 2011–2016 Real-Time Systems Lab, TU Darmstadt. Anthony Anjorin, Erika Burdon, Frederik Deckwerth, Roland Kluge, Lars Kliegel, Marius Lauder, Erhan Leblebici, Daniel Tögel, David Marx, Lars Patzina, Sven Patzina, Alexander Schleich, Sascha Edwin Zander, Jerome Reinländer, Martin Wieber, and contributors. All rights reserved.

This document is free; you can redistribute it and/or modify it under the terms of the GNU Free Documentation License as published by the Free Software Foundation; either version 1.3 of the License, or (at your option) any later version. Please visit <http://www.gnu.org/copyleft/fdl.html> to find the full text of the license.

For further information contact us at contact@emoflon.org.

The eMoflon team

Darmstadt, Germany (May 2016)

Contents

1	A language definition problem?	2
2	Abstract syntax and static semantics	5
3	Creating instances	32
4	eMoflon’s graph viewer	36
5	Introduction to injections	39
6	Leitner’s Box GUI	44
7	Conclusion and next steps	46
	Glossary	47

Part II:

Leitner's Learning Box

URL of this document: <https://emoflon.github.io/eclipse-plugin/beta/handbook/part2.pdf>

The toughest part of learning a new language is often building up a sufficient vocabulary. This is usually accomplished by repeating a long list of words again and again until they stick. A *Leitner's learning box*¹ is a simple but ingenious little contraption to support this tedious process of memorization.

As depicted in Figure 0.1, it consists of a series of compartments or partitions usually of increasing size. The content to be memorized is written on a series of cards which are initially placed in the first partition. All cards in the first partition should be repeated everyday and cards that have been successfully memorized are placed in the next partition. Cards in all other partitions are only repeated when the corresponding partition is full and cards that are answered correctly are moved one partition forward in the box. Challenging cards that have been forgotten are treated as brand new cards and are always returned to the first partition, regardless of how far in the box they have progressed.

These “rules” are depicted by the green and red arrows in Figure 0.1. The basic idea is to repeat difficult cards as often as necessary and not waste time on easy cards which are only repeated now and then to keep them in memory. The increasing size of the partitions represent how words are easily placed in our limited short term memory and slowly move in our theoretically unlimited long term memory if practised often enough.

¹http://en.wikipedia.org/wiki/Leitner_system




Figure 0.1: Possible *concrete syntax* of a Leitner's learning box

A learning box is an interesting system, because it consists clearly of a static structure (the box, partitions with varying sizes, and cards with two sides of content) and a set of rules that describe the dynamic aspects (behaviour) of the system. In this Part II of the eMoflon handbook, we shall build a complete learning box from scratch in a model-driven fashion and use it to introduce fundamental concepts in metamodeling and *Model-Driven Software Development* (MDSD) in general. We'll begin with a short discussion, then develop the box's abstract syntax, learn about eMoflon's validation, create some dynamic instances of the box, use injections to integrate a hand-written implementation of a method with generated code from the model, and get to know a small GUI that will allow us to take the implemented learning box for a spin.

1 A language definition problem?

As in any area of study, metamodeling has its fair share of buzz words used by experts to communicate concisely. Although some concepts might seem quite abstract for a beginner, a well defined vocabulary is important so we know exactly what we are talking about.

The first step is understanding that metamodeling equates to language definition. This means that the task of building a system like our learning box can be viewed as defining a suitable language that can be used to describe

such systems. This language oriented approach has a lot of advantages including a natural support for product lines (individual products are valid members of the language) and a clear separation between platform independent and platform specific details.

So what constitutes a language? The first question is obviously what the building blocks of your language “look” like. Will your language be textual? Visual? This representation is referred to as the *concrete syntax* of a language and is basically an interface to end users who use the language. In the case of our learning box, Figure 0.1 can be viewed as a possible concrete syntax. As we are building a learning box as a software system however, our actual concrete syntax will be composed of GUI elements like buttons, drop-down menus and text fields.

*Concrete
Syntax*

Irrespective of what a language looks like, members of the language must adhere to the same set of “rules”. For a natural language like English, this set of rules is usually called a *grammar*. In metamodeling, however, everything is represented as a graph of some kind and, although the concept of a *graph grammar* is also quite well-spread and understood, metamodelers often use a *type graph* that declaratively defines what types and relations constitute a language.

*Grammar
Graph
Grammar
Type Graph*

A graph that is a member of your language must *conform* to the corresponding type graph for the language. To be more precise, it must be possible to type the graph according to the type graph - the types and relations used in the graph must exist in the type graph and not contradict the structure defined there. This way of defining membership to a language has many parallels to the class-object relationship in the object-oriented (OO) paradigm and should seem very familiar for any programmer used to OO. This type graph is referred to as the *abstract syntax* of a language.

*Abstract
Syntax*

Often, one might want to further constrain a language, beyond simple typing rules. This can be accomplished with a further set of rules or constraints that members of the language must fulfil in addition to being conform to the type graph. These further constraints are referred to as the *static semantics* of a language.

*Static
Semantics*

With these few basic concepts, we can now introduce a further and central concept in metamodeling, the *metamodel*, which is basically a simple class diagram. A metamodel defines not only the abstract syntax of a language but also some basic constraints (a part of the static semantics).

Metamodel

In analogy to the “everything is an object” principle in the OO paradigm, in metamodeling, everything is a model! This principle is called *unification*, and has many advantages. If everything is a model, a metamodel that defines (at least a part of) a language must be a model itself. This means that it con-

*Unification
Modelling
Language*

forms to some *meta-metamodel* which in turn defines a *(meta)modelling language* or *meta-language*. For metamodeling with eMoflon, we support *Ecore* as a modelling language and it defines types like **EClass** and **EReference**, which we will be using to specify our metamodels. Alternate modelling languages include MOF, UML and Kermeta.

*Meta-metamodel
Meta-Language*

Thinking back to our learning box, we can define the types and relations we want to allow. We want an entire box of flashcards where each card is contained within a partition, and each partition is contained within the box. Multiplicities are an example of static semantics that do not belong to the abstract syntax, but can nonetheless be expressed in a metamodel. An example could be that a card can only ever exist in one partition, or that a partition can have either one **next** partition, or none at all.

More complex constraints that cannot be expressed in a metamodel are usually specified using an extra *constraint language* such as the Object Constraint Language (OCL). This idea, however, is beyond the scope of this handbook. We'll stick to metamodels without an extra constraint language.

Constraint Language

In addition to its static structure, every system has certain dynamic aspects that describe the system's behaviour and how it reacts to external stimulus or evolves over time. In a language, these rules that govern the dynamic behaviour of a system are referred to collectively as the *dynamic semantics* of the language. Although these rules can be defined as a set of separate *model transformations*, we take a holistic approach and advocate integrating the transformations directly in the metamodel as operations. This naturally fits quite nicely into the OO paradigm.

Dynamic Semantics

A short recap: We have learned that metamodeling starts with defining a suitable language. For the moment, we know that a language comprises a concrete syntax (how the language “looks”), an abstract syntax (types and relations of the underlying graph structure), and static semantics (further constraints that members of the language must fulfil). Metamodels are used to define the abstract syntax, and a part of the static semantics of a language, while *models* are graphs that conform to some metamodel (i.e., can be typed according to the abstract syntax and must adhere to the static semantics).

Model

This handbook is meant to be hands-on, so enough theory! Lets define, step-by-step, a metamodel for a learning box using our tool, eMoflon.

2 Abstract syntax and static semantics

The first step in creating any metamodel is defining the abstract syntax, also known as the type graph. This involves defining each class, its attributes, references, and method signatures.

If you completed the demo in Part I, your Eclipse workspace will look slightly different than ours depicted in the screenshots. In an effort to keep things as clear as possible, we have removed those files from our package explorer, but still recommend keeping them for future reference.

Additionally, if you're continuing from the demo, you can begin modelling this project in two different ways. You can either develop your metamodel in the same workspace as the demo, or create a new one. Either way, please note that the steps are exactly the same, but our project browsers in EA may not exactly match. This handbook has assumed you prefer the latter.

2.1 Getting started in EA

- To begin, navigate to “New Metamodel Project” (Figure 2.1) and start a new project named `LeitnersLearningBox` (Figure 2.2). Open the empty `.eap` file in EA.



Figure 2.1: “New Metamodel Project” button

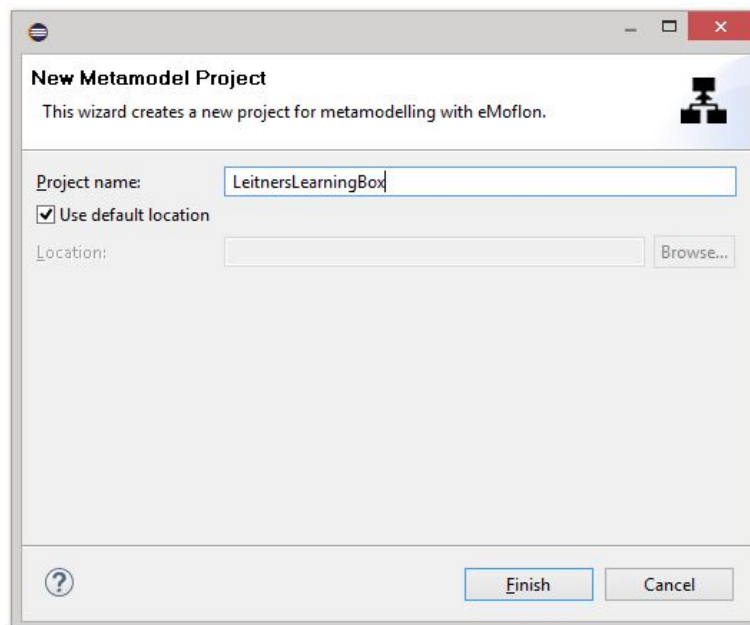


Figure 2.2: Starting a new visual project

- In EA, select your working set and press the “Add a Package” button (Figure 2.3).




Figure 2.3: Add a new package to MyWorkingSet

- In the dialogue that pops up (Figure 2.4), enter `LearningBoxLanguage` as the name of the new package. In this case select **Package Only** and click **OK**. Later you can select **Create Diagram** to skip the next step.

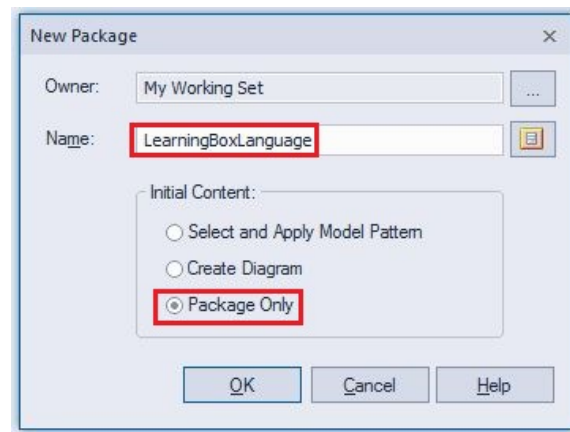


Figure 2.4: Enter the name of the new package

- Your **Project Browser** should now resemble Figure 2.5.




Figure 2.5: State after creating the new package

- Now select your new package and create a “New Diagram” (Figure 2.6).




Figure 2.6: Add a diagram

- In the dialogue that appears (Figure 2.7), choose **eMoflon Ecore Diagrams** and press **OK**.

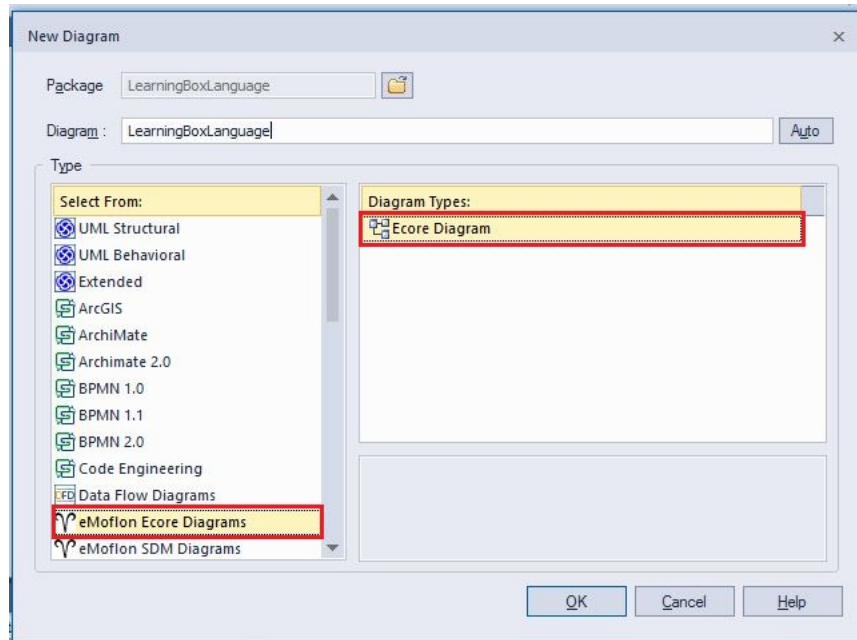


Figure 2.7: Select the ecore diagram type

- After creating the new diagram, your **Project Browser** should now resemble Figure 2.8. You'll notice that your **LearningBoxLanguage** package has been transformed into an **EPackage**.




Figure 2.8: State after creating diagram

- You can now already export your project to Eclipse,² then refresh your **Package Explorer**. A new node, **My Working Set**³ should have appeared containing your **EPackage** (Figure 2.9). You can see that a **LearningBoxLanguage.ecore** file has been generated, and placed in “model.” This is your metamodel that will contain all future types you create in your diagrams.

²If unsure how to perform this step, please refer to Part I, Section 2.1

³If you do not have the two distinct nodes, ensure your “Top Level Elements” are set to **Working Sets**

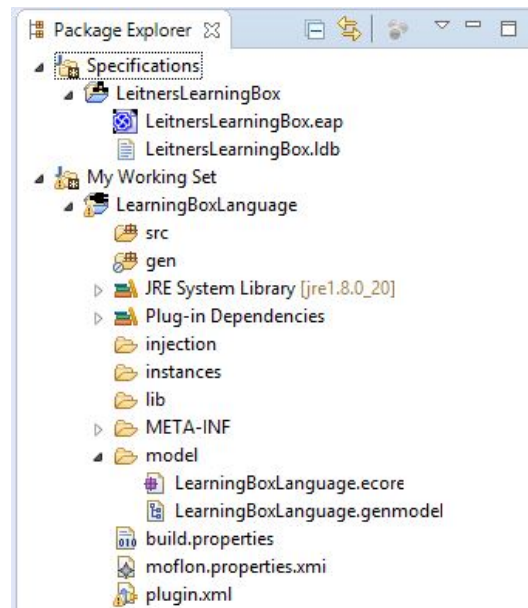


Figure 2.9: Initial export to Eclipse

- If you're interested in reviewing the overall project structure and the purposes of certain files and folders, read Section 4.1 from Part I. Otherwise, continue to the next section to learn how to declare classes and attributes.

2.2 Declaring classes and attributes

- Return to EA, and double-click your `LearningBoxLanguage` diagram to ensure it's open.
- There are two ways for you to create your first `EClass`. First, to the left of the workbench, a *Toolbox* containing the Ecore types available for metamodeling should have appeared (Figure 2.10).⁴ Click on the `EClass` icon then somewhere in the diagram to create a new object. Alternatively, you can click in the diagram and press `space` to invoke the toolbox context menu, then select `EClass`.

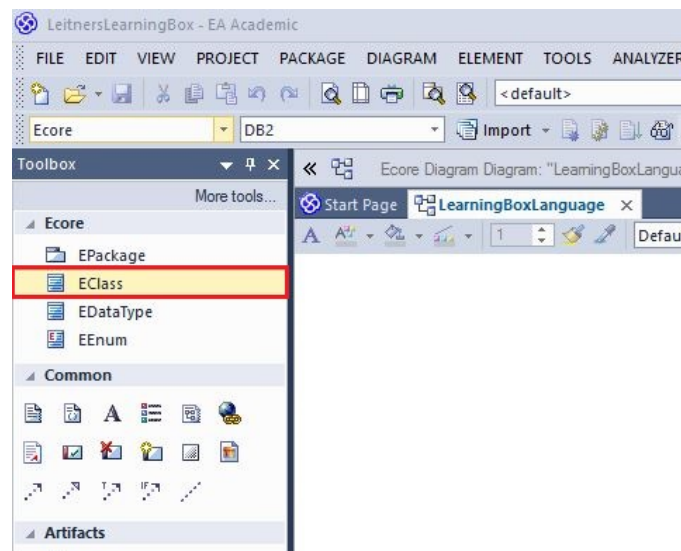


Figure 2.10: Create an EClass

- In the dialogue that pops-up, set `Box` as the name and click `OK` (Figure 2.11). This dialogue can always be invoked again by double-clicking the `EClass`, or by pressing `Alt` and single-clicking. It contains many other properties that we'll investigate later in the handbook. In general, a similar properties dialogue can be opened in the same fashion for almost every element in EA.

⁴If not, choose "Diagram/Diagram Toolbox" to show the current toolbox (`Alt+ 5`)

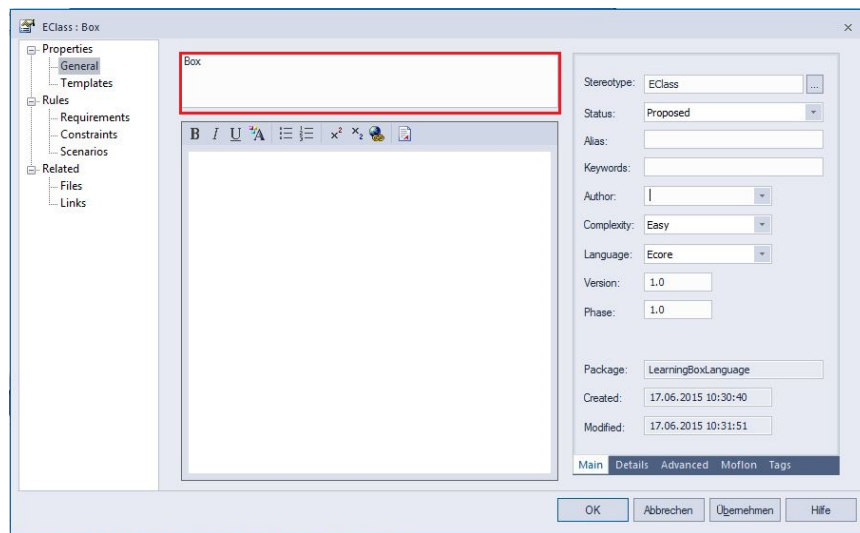


Figure 2.11: Edit the properties of an EClass

- After creating Box, your EA workspace should resemble Figure 2.12.

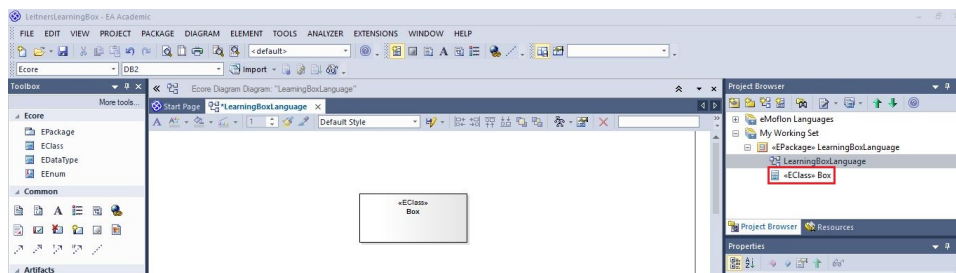


Figure 2.12: State after creating Box

- Now create the **Partition** and **Card** EClasses the same way, until your workspace resembles Figure 2.13. These are the main classes of your learning box metamodel.
- Lets add some attributes! Either right-click on **Box** to activate the context menu and choose “Features & Properties/Attributes..” (Figure 2.14), or press F9 to open the editing dialogue.

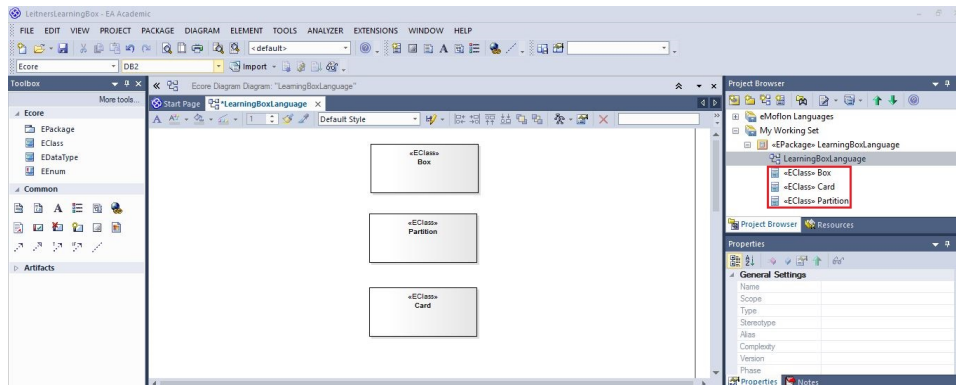


Figure 2.13: All EClasses for the metamodel

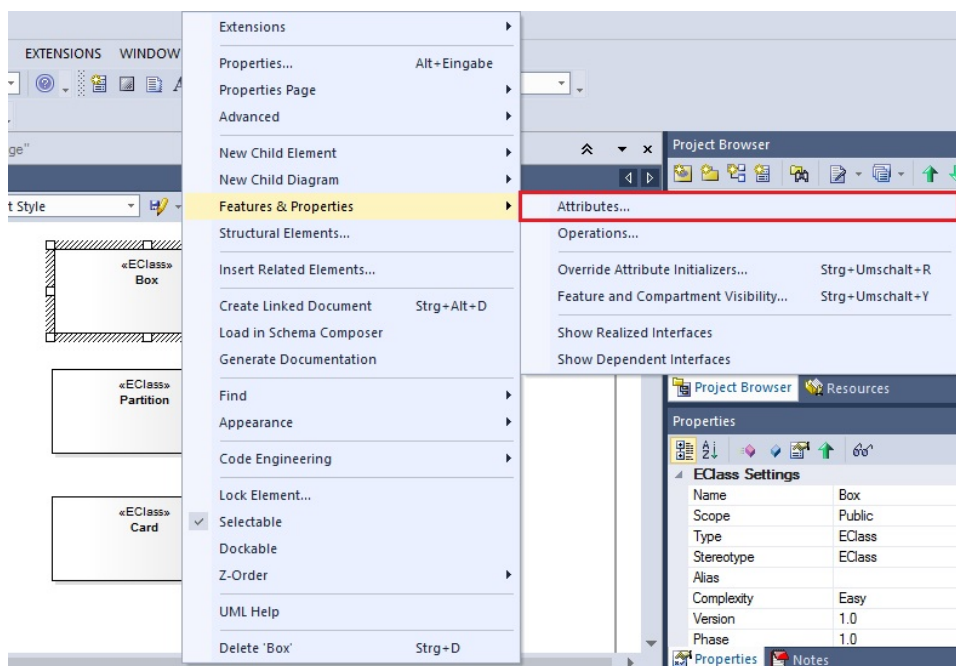


Figure 2.14: Context menu for an EClass

- Enter **name** as the name of the attribute, select **EString** as its type from the drop-down menu, and press **Close** (Figure 2.15). New attributes for the same EClass can be added by clicking on **New Attribute...**

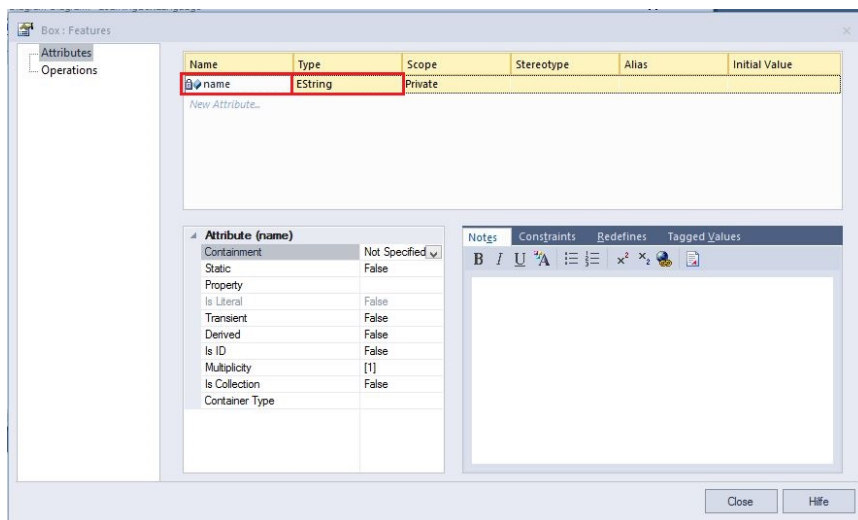


Figure 2.15: Adding attributes to an EClass

- Add the remaining attributes analogously to each EClass until your workspace resembles Figure 2.16.
- Save and export to Eclipse. After refreshing your workspace, your **.ecore** model can now be expanded as it includes every class and attribute from your metamodel. So far, so good!

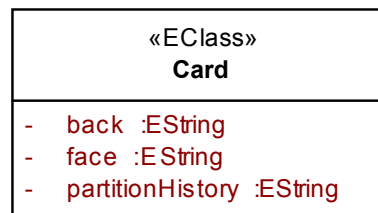




Figure 2.16: Main EClasses declared with their attributes

2.3 Connecting your classes

At this point, you’ve declared your types and their attributes, but what good are those if you can’t connect them to each other? We need to create some *EReferences*!

EReference

There are four properties that must be set in order to create an *EReference*: the target and source *Role*, *Navigability*, *Multiplicity* and *Aggregation*, all of which are declared differently in each syntax, but let’s first review the concepts since they’re basically the same.

A *Role* clarifies which way a reference is ‘pointing.’ In other words, the *Role source* and *target* roles determine the direction of the connection.

Navigable ends are mapped to class attributes with getters and setters in Java, and therefore *must* have a specified name and multiplicity for successful code generation. Corresponding values for *Non-Navigable* ends can be regarded as additional documentation, and do not have to be specified.

Navigability

The *Multiplicity* of a reference controls if the relation is mapped to a (Java) collection (*, ‘1..*’, ‘0..*’), or to a single valued class attribute (‘1’, ‘0..1’). We’ll explain this setting in detail later.

Multiplicity

The *Aggregation* values of a reference can be either none or composite. Composite means that the current role is that of a *container* for the opposite role. You’ll see in our example that **Box** is a container for several **Partitions**. This has a series of consequences: (1) every element must have a container, (2) an element cannot be in more than one container at the same time, and (3) a container’s contents are deleted together with the container. Conversely, non-composite (denoted by “none”) means that the current role is not that of a container, and the rules for containment do not hold (in other words, the reference is a simple ‘pointer’).

Aggregation
Container

Creating EReferences in EA

- A fundamental gesture in EA is *Quick Link*. Quick Link is used to create EReferences between elements in a context-sensitive manner. To use quick link, choose an element and note the little black arrow in its top-right corner (Figure 2.17).




Figure 2.17: Quick Link is a central gesture in EA

- Click this black arrow and ‘pull’ to the element you wish to link to. To start, quick-link from **Box** to **Partition**. In the context menu that appears, select “Create Bidirectional EReference” (Figure 2.18).

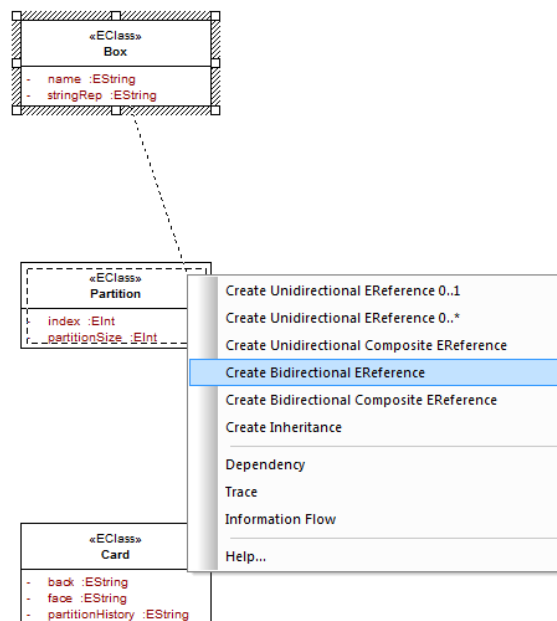


Figure 2.18: Create an EReference via Quick Link

- Double click the EReference to invoke a dialogue. In this window you can adjust all relevant settings. Feel free to leave the **Name** value blank - this property is only used for documentation purposes, and is not relevant for code generation.

- Within this dialogue, go to “Role(s),” and compare the relevant values in Figure 2.19 for the *source* end of the EReference (the **Box** role). As you can see, the default source is set to the EClass you linked from, while the default target is the EClass you linked to. In this window, do not forget to confirm and modify the **Role**, **Navigability**, **Multiplicity**, and **Aggregation** settings for the source as required. Repeat the process for the **Target Role**.

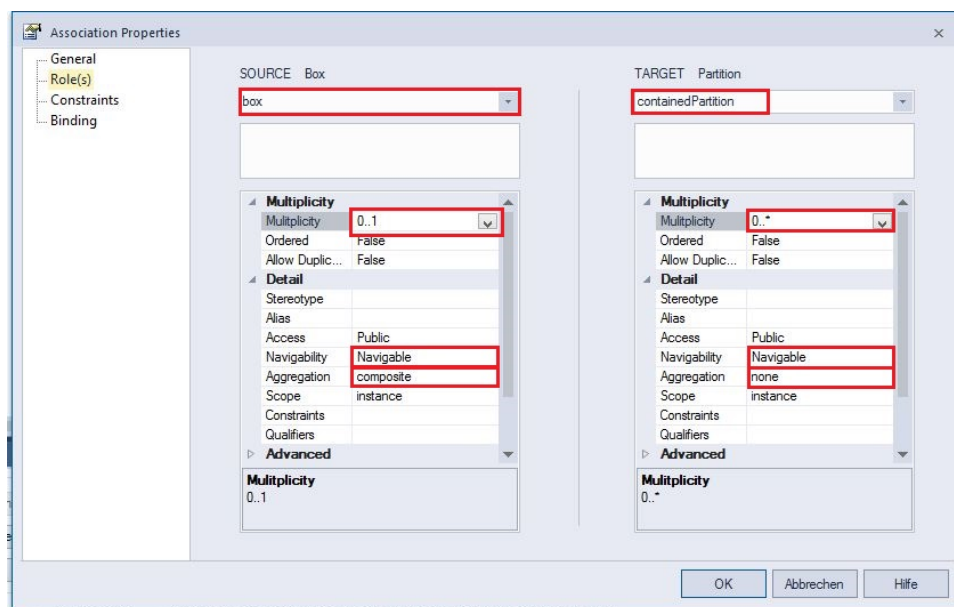


Figure 2.19: Properties for the source and target role of an EReference

To review these properties, the first value you edited was the role name. The **Navigability** value should have been automatically set to **Navigable**. Without these two settings, getter and setter methods will not be generated.

Next, you set the **Multiplicity** value. In your source role (**Box**), you have allowed the creation of up to one target (**Partition**) reference for every connected source (**box**). This means you could not have a single target connected to two sources (i.e., one partition that belongs to two boxes). In the target (**Partition**) role, you have specified that any source (in our case, **box**) can have any positive-sized number of targets. Figure 2.20 sketches this schematically.




Figure 2.20: The target and source roles of Leitner's Learning Box

Finally, you set the **Aggregation** value. In this case, **box** is a container for **Partitions**, and **containedPartition** is consequently not.

- Take a moment to review how the **Aggregation** settings extend the **Multiplicity** rules. If you've done everything right, your metamodel should now resemble Figure 2.21, with a single *bidirectional EReference* between **Box** and **Partition**.

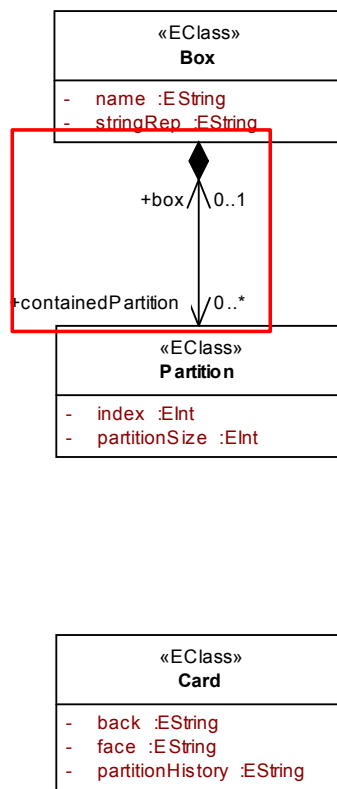


Figure 2.21: Box contains Partitions

- Following the same process, create two unidirectional self-EReferences for **Partition** (for next and previous **Partitions**), and a second bidirectional composite EReference⁵ between **Partition** and **Card** (Figure 2.22).

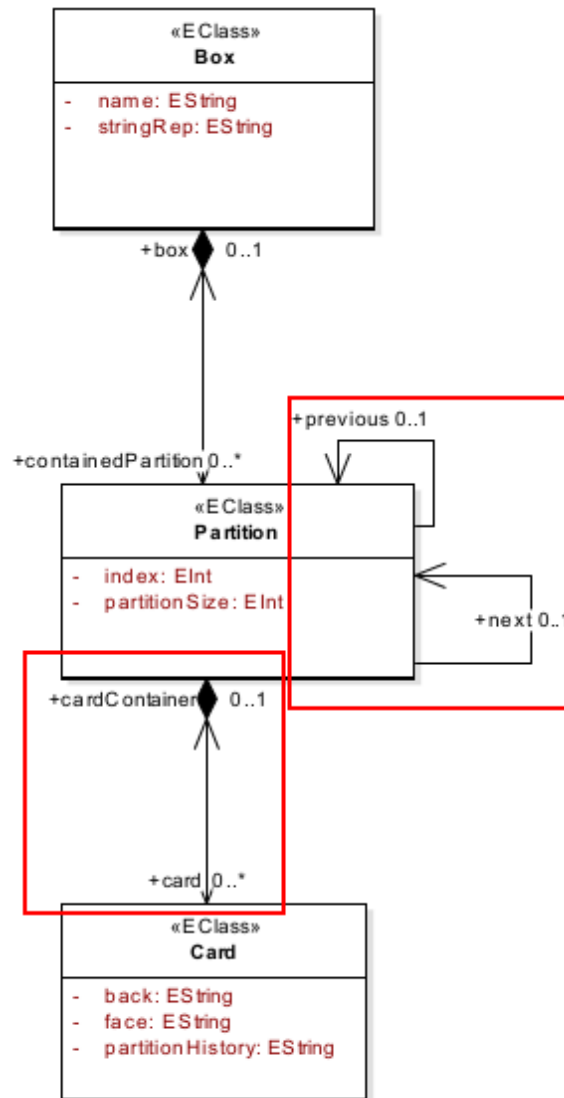


Figure 2.22: All relations in our metamodel

⁵To be precise, *all* EReferences in Ecore are actually unidirectional. A “bidirectional” EReference in our metamodel is really two mapped EReferences that are opposites of each other. We however, believe it is simpler to handle these pairs as single EReferences, and prefer this concise concrete syntax.

- You'll notice that the connection between **Card** and **Partition** is similar to that between **Partition** and **Box**. This makes sense as a partition should be able to hold an unlimited amount of cards, but a card can only belong to one partition at a time.
- Export your diagram to Eclipse and refresh your workspace. Your Ecore metamodel file in “model/LearningBoxLanguage.ecore” should now resemble Figure 2.23.

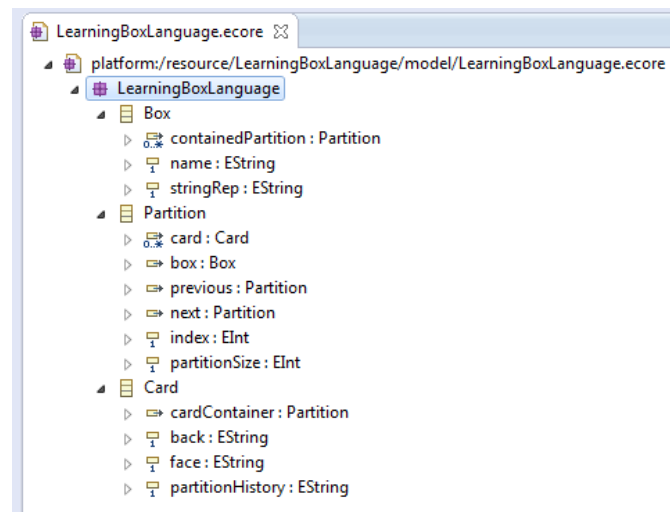


Figure 2.23: Refreshed Ecore file with all EReferences

- All the required attributes and references for your learning box have now been set up.

2.4 Method Signatures

To finish defining our types, let's define the *signatures* of some operations that they'll eventually support. *Operation Signature*

- Select **Partition** and either right-click to invoke the context-menu (Figure 2.24) and choose “Features & Properties/Operations..” or simply press F10.

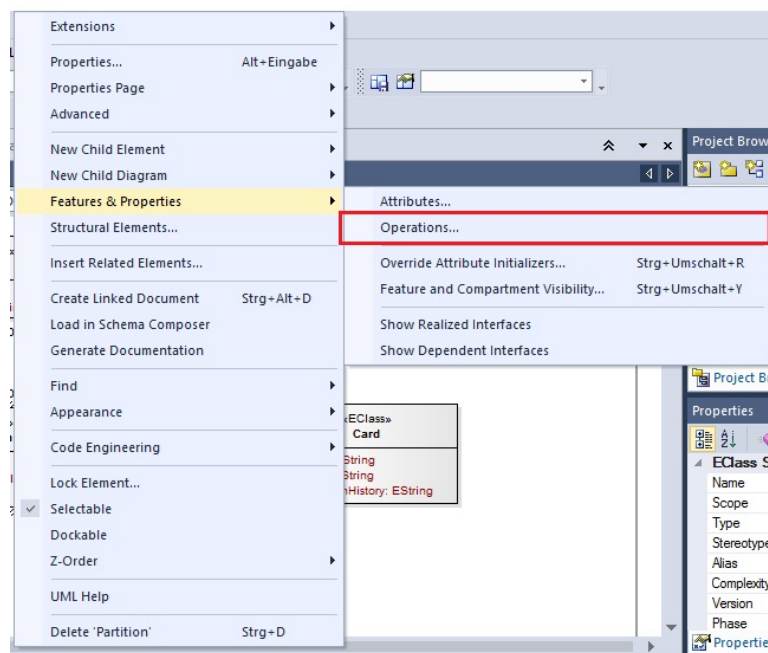


Figure 2.24: Add an operation

- In the dialogue that pops-up (Figure 2.25), enter **empty** as the Name of the operation and **void** as the Return Type.
- In the same dialogue, click on **New Operation...** to add a second operation, **removeCard**, and edit the values as seen in Figure 2.26. Notice that the Return Type can be chosen by either the drop-down menu, or via direct typing. For types you've established in the metamodel (e.g. **Card**) you have to use 'Select Type...' from the drop-down menu.

Very important: Non-primitive types *must* be chosen via 'Select Type...' in the drop-down menu. It allows you to browse for the corresponding elements in your project. Simply typing them won't work!

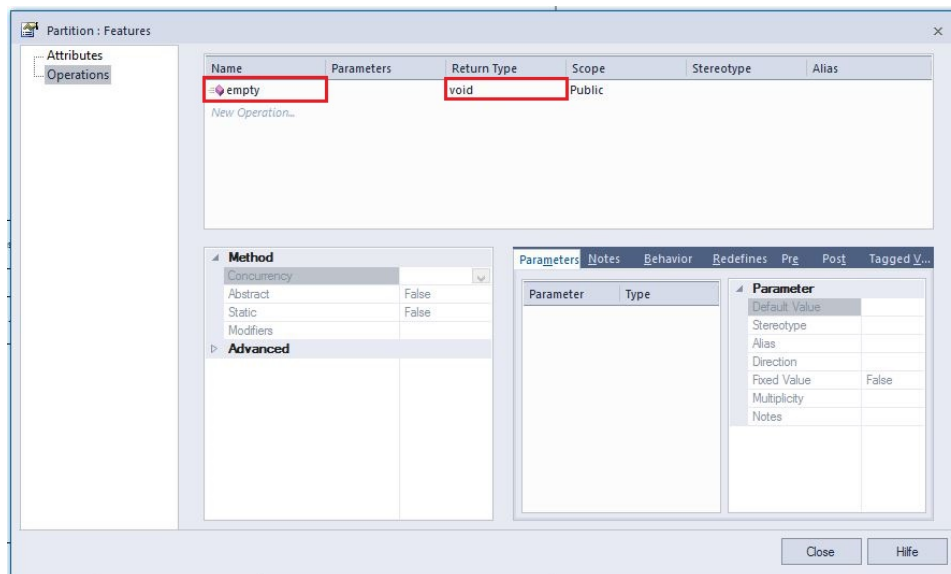


Figure 2.25: EClass properties editor

- ▶ Parameters can be added by selecting **Parameters** and completing the dialogue (Figure 2.26). Please remember that you must also use either the drop-down menu, or direct typing to select the type or else validation will fail.
- ▶ Repeat this process for the **check** operation (with the two parameters **card:Card**, **guess:EString**) that returns an **EBoolean**.
- ▶ If you've done everything right, your dialogue should now contain three methods - **check**, **empty**, and **removeCard** - each with the corresponding parameters and return types in Figure 2.27.
- ▶ Add all operations analogously for **Box** and **Card** until your metamodel closely resembles Figure 2.28.⁶
- ▶ To finish, export the metamodel for code generation in Eclipse, and examine the model once again. Each signature should have appeared in their respective EClass.

⁶Please note that names of parameters may not be displayed by default in EA

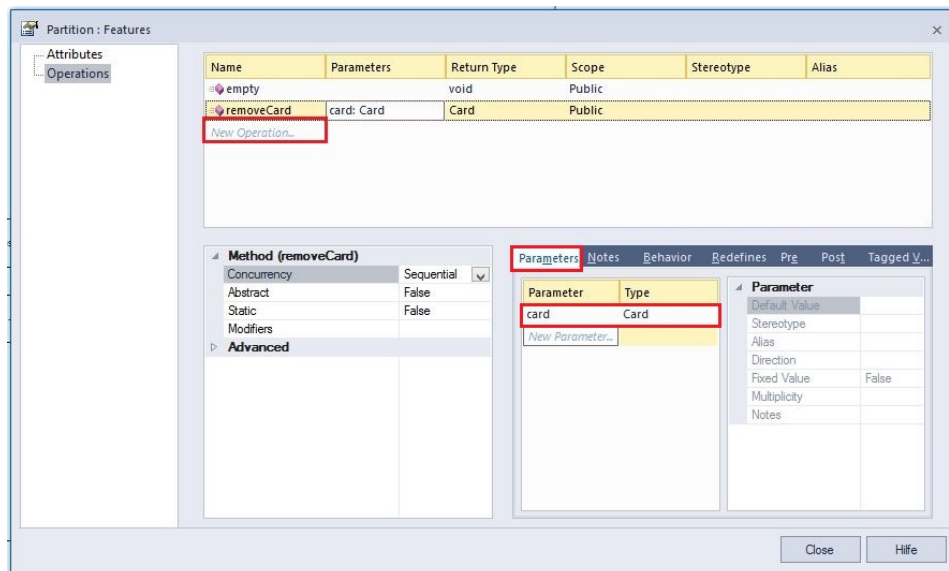


Figure 2.26: Parameters and return type

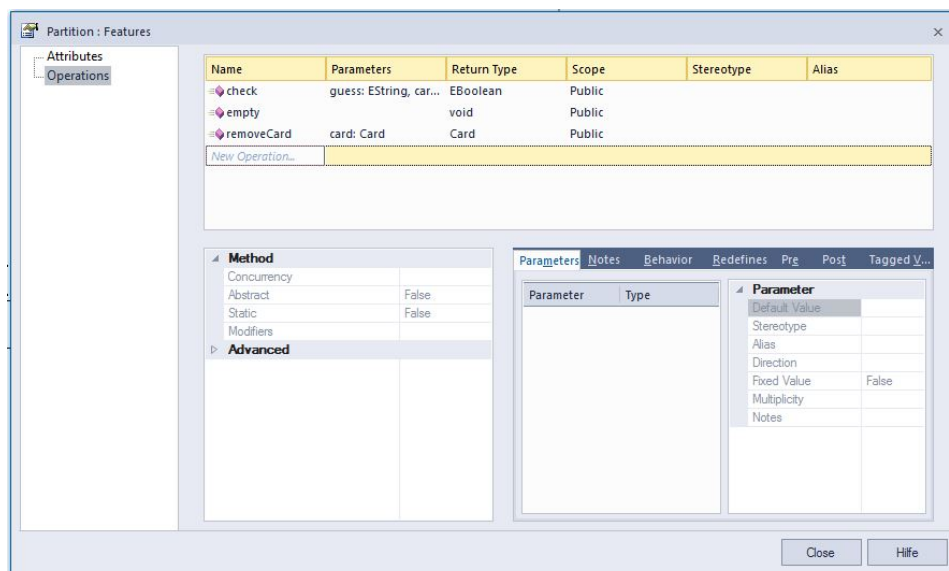


Figure 2.27: All operations in Partition

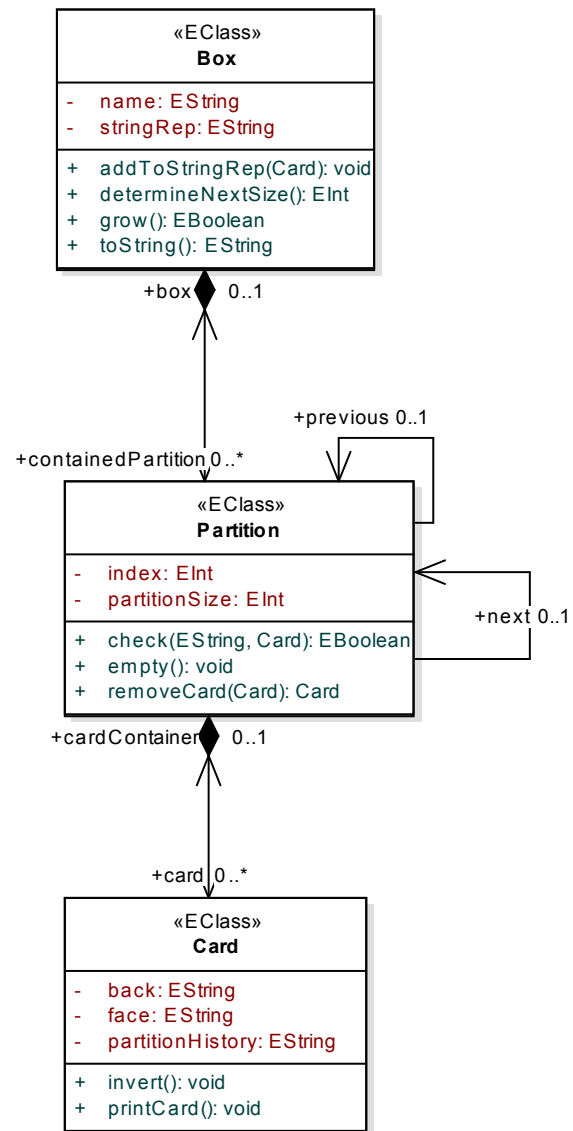


Figure 2.28: Complete metamodel for our learning box

2.5 eMoflon validation support in EA

Our EA extension provides rudimentary support for validating your meta-model. Validation results are displayed and, in some cases, even “quick fixes” to automatically solve the problems are offered. In addition to reviewing your model, the validation option automatically exports the current model to your eclipse workspace if no problems were detected.

- If not already active, make the eMoflon control panel visible in EA by choosing “Extensions/Add-In Windows”. This should display a new output window, as depicted in Figure 2.29. Many users prefer this interface as it provides quick access to all of eMoflon’s features, as opposed to the drop down menu under “Extensions/MOFLON::Ecore Addin” which only offers limited functionality.

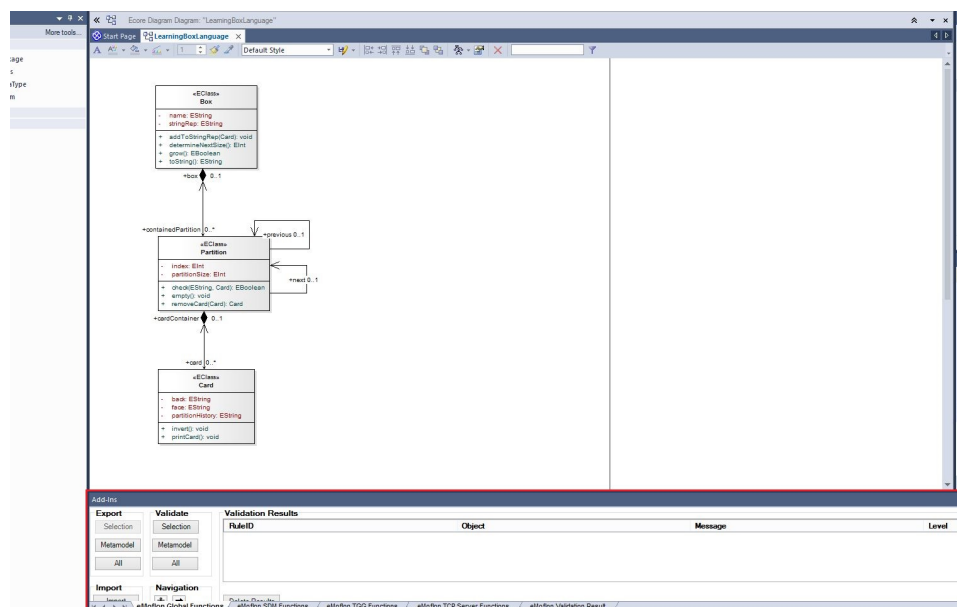


Figure 2.29: Activating the validation output window

- To start the validation, choose “Validate all” in the “Validate” section of the control panel (Figure 2.30). If you haven’t made any mistakes while modelling your `LearningBoxLanguage` so far, the validation results window should remain empty, indicating your metamodels are free of errors.

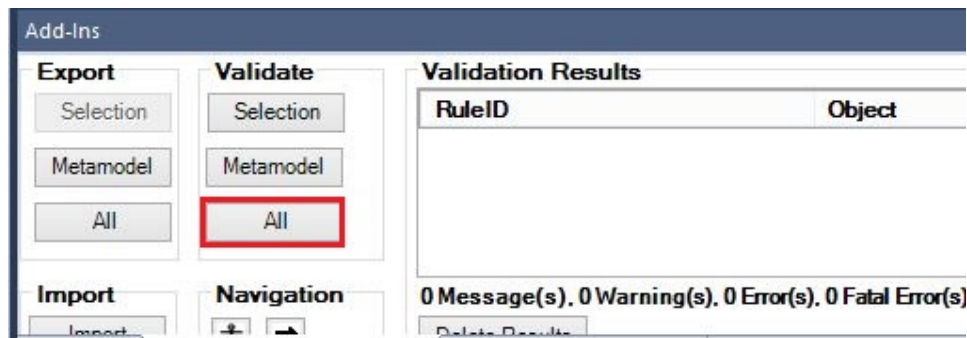


Figure 2.30: Starting the validation

If an error did appear, the validation system attempts to suggest a “Quick Fix.” Why don’t we examine the validation and quick fix features in detail? Let’s add two small modelling errors in `LearningBoxLanguage`.

- Create a new EClass in the `LearningBoxLanguage` diagram. You can retain the default name `EClass1`. Let’s assume you wish to delete this class from your metamodel.
- Select the rouge class in the diagram and press the **Delete** button on your keyboard. Note that this only deleted it from the current diagram and `EClass1` still exists in the project browser (and thus in your metamodel).
- Run the validation test, and notice the new **Information** message in the validation output (Figure 2.31).

Figure 2.31: Validation information error: element still exists

This message informs you that `EClass1` is not on any diagram, and seeing as it is still in the metamodel, that this *could* be a mistake. As you can see, just pressing the **Delete** button is not the proper way of removing an EClass from a metamodel - It only removes it from the current diagram!⁷

⁷Deleting elements properly and other EA specific aspects are discussed in detail in Part VI: Miscellaneous

- Suppose you were inspecting a different diagram, and were not on the current screen. To navigate to the problematic element in the **Project Browser**, click *once* on the information message.
- To check to see if there are any quick fixes available, *double-click* the information message to invoke the “QuickFix” dialogue. In this case, there are two potential solutions - add the element to the current diagram or (properly) delete the element from the metamodel (Figure 2.32). Since the latter was the original intent, click **Ok**.

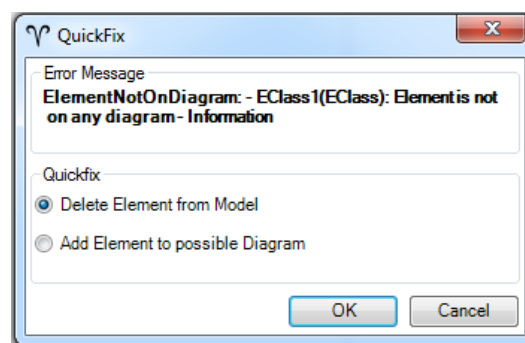


Figure 2.32: Quick fix for elements that are not on any diagram

- **EClass1** should now be correctly removed from your metamodel. Your metamodel should now be error-free again as indicated by the validation output window.
- To make an error that leads to a more critical message than “information,” double-click the navigable **EReference** end **previous** of the **EClass** **Partition**, and delete its role name as depicted in Figure 2.33. Affirm with **OK**.
- You should now see a new **Fatal Error** in the validation output, stating that a navigable end *must* have a role name. Close all windows, then single click on the error once to open the relevant diagram and highlight the invalid element on the diagram. Double click the error to view the quick fix menu (Figure 2.34). As navigable references are mapped to data members in a Java class, omitting the name of a navigable reference makes code generation impossible (data members ,i.e., class variables, must have a name).
- Given there are no automatic solutions, correct your metamodel manually by setting the name of the navigable **EReference** back to **previous**.

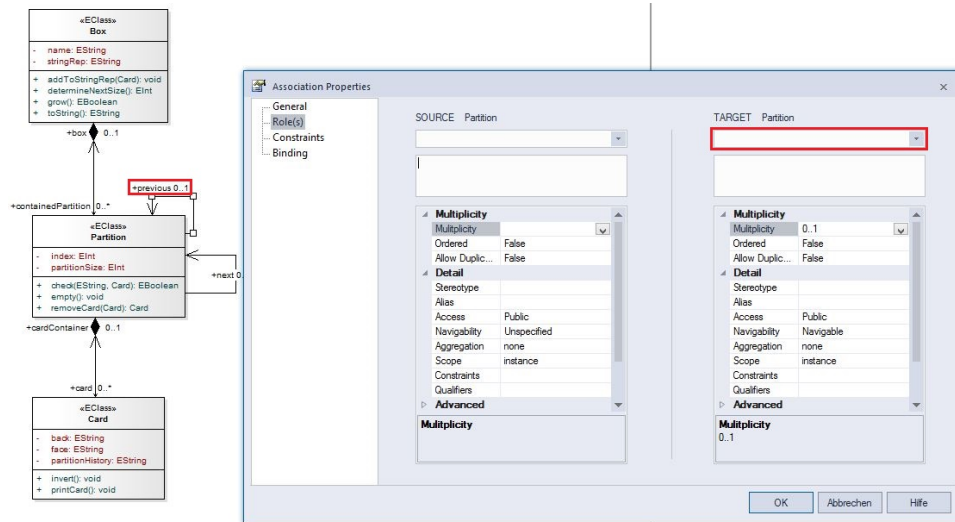


Figure 2.33: Deleting a navigable role name of an EReference




Figure 2.34: Fatal error after deleting a navigable role name

-
- Ensure that your metamodel closely resembles Figure 2.28 again, and that there are no error messages before proceeding.

As you may have noticed, eMoflon distinguishes between five different types of validation messages:

Information:

This is only a hint for the user and can be safely ignored if you know what you're doing. Export and code generation should be possible, but certain naming/modelling conventions are violated, or a problematic situation has been detected.

Warning:

Export and code generation is possible, but only with defaults and automatic corrections applied by the code generator. As this might not be what the user wants, such cases are flagged as warnings (e.g., omitting the multiplicity at references which is automatically set by the code generator to 1). Being as explicit as possible is often better than relying on defaults.

Error:

Although the metamodel can be exported from EA, it is not Ecore conform, and code generation will not be possible.

Fatal Error:

The metamodel cannot be exported as required information, such as names or classifiers of model elements is incorrectly set or missing.

Eclipse Error:

Display error messages produced by our Eclipse plugin after an unsuccessful attempt to generate code.

2.6 Reviewing your metamodel

Before moving on, let's take a step back and review what we have accomplished. We have modelled a **Box** that can contain an arbitrary amount of **Partitions**. A **Partition** in the **Box** has a **next** and **previous Partition** that can be set. Finally, **Partitions** contain **Cards**.

A **Box** has a **name**, and can be extended by calling **grow**. A **Box** can print out its contents via the **toString** method.

The main method of the learning box is **Partition::check**, which takes a **Card** and the user's **EString** guess, and returns a **true** or **false** value.

A **Partition** can also **remove** a specific **Card**, or empty itself of all existing **Cards**. Last but not least, a **Partition** has a **partitionSize** to indicate how many cards it should hold. Too many cards in the first partition could indicate that not enough time has been dedicated to learning the terms. Too many near the end of the box could show that the vocabulary set is too easy, and probably already mastered.

A **Card** contains the actual content to be learned as a question on the card's **face** and the answer on the card's **back**. **Cards** also maintain a **partition-History**, which can be used to keep track of how often a **Card** has been answered incorrectly. This may indicate how difficult the **Card** is for a specific user, and remind them to spend more time on it. When learning a language, it makes sense to be able to swap the target and source language and this is supported by **Card** via **invert** (turns the card around).

Examine the generated files in "gen", especially the default implementation for those methods that currently just throw an **OperationNotSupported** exception. We shall see in later parts of this handbook how our code generator supports injecting hand-written implementation of methods into generated methods and classes. With eMoflon however, we can actually model a large part of the dynamic semantics with ease, and only need to implement small helper methods (such as those for string manipulation) by hand.

If you have had problems with this section, and, despite firmly believing everything is correct, things *still* don't work, feel free to contact us at: contact@emoflon.org.

3 Creating instances

Before diving into modelling dynamic behaviour in Part III, let's have a brief look at how to create a concrete *instance* of your metamodel in Eclipse.

In the following section, we use *metamodel* and *instance* (model) to differentiate between models that represent the abstract syntax and static semantics of a domain specific language (metamodel), and those that are expressed *in* such a language (instance models of the metamodel).

- ▶ To create an instance model, navigate to the generated `model` folder in your `LearningBoxLanguage` project. Double-click the `LearningBoxLanguage.ecore` model to invoke the *Ecore model editor*.
- ▶ To create a concrete instance of the metamodel, you must select an EClass that will become the root element of the new instance. For our example, right-click `Box`, and navigate to “Create Dynamic Instance” from the context menu, as depicted in Figure 3.1.
- ▶ A dialogue should appear asking where the instance model file should be persisted. Save your instances according to convention in a folder named “instances,” which is automatically created in every new repository project. Last but not least, enter `Box.xmi` as the name of the instance model (Figure 3.2).
- ▶ Press **Finish**, and a generic model editor should open for your new instance model. This editor works just like the previous Ecore model editor except it's “generic,” meaning it allows you to create and edit an instance of *any* metamodel, not only instances of Ecore (i.e., metamodels).

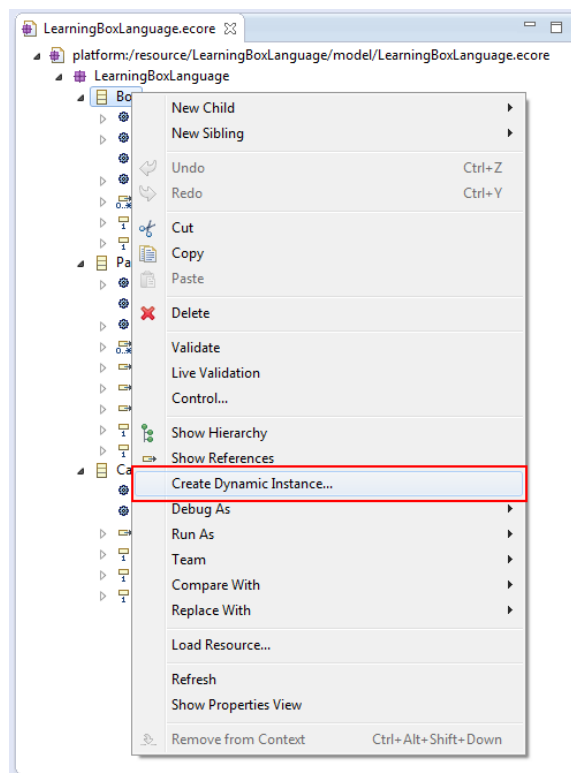


Figure 3.1: Context menu of an EClass in the Ecore editor

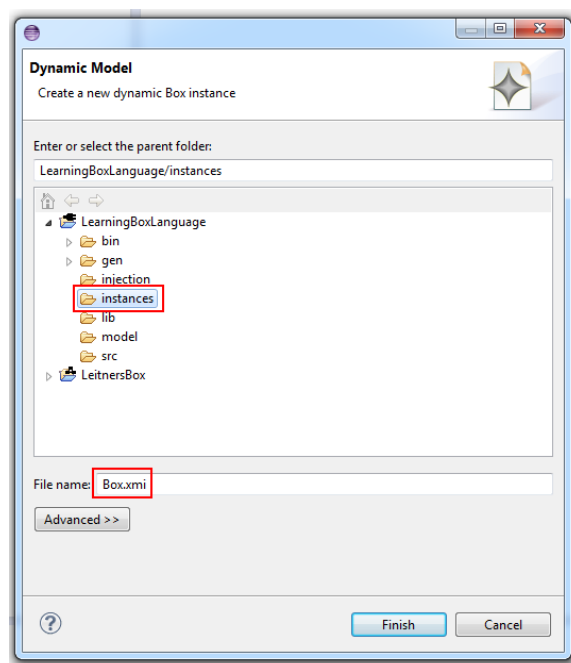


Figure 3.2: Dialogue for creating a dynamic model instance

- You can populate your instance by adding new children or siblings via a right-click of an element to invoke the context-menu depicted in Figure 3.3. Note that EMF supports you by respecting your metamodel, and reducing the choice of available elements to valid types only.⁸




Figure 3.3: Context menu for creating model elements

- Let's try building a vocabulary set. Fill your box with three partitions, with two cards in each. Save your model by pressing **Ctrl+S** and confirm the save by closing it, then reloading via a simple double-click.
- Double-click on one of the partitions to bring up the "Properties" tab in the window below the editor (Figure 3.4). Here you'll see the attributes you defined earlier in each EClass. The **Box**, **Next**, and **Previous** values are its (undefined) EReferences,⁹ while **Index** is a partition's unique identification value, and **Partition Size** is the recommended number of cards the partition should contain before being tested.

⁸This depends on the current context. Try it out!

⁹Please note that the editor tab capitalizes the first letter of each property

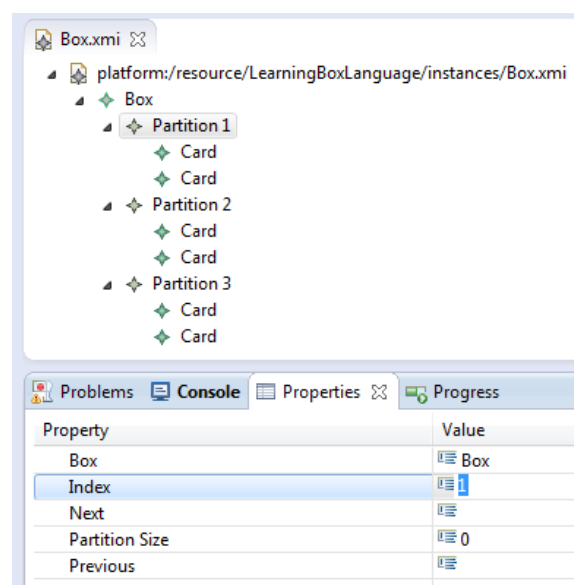


Figure 3.4: Change the **Index** value in the partition's property tab

- Pick a number - any one you like - and update each partition's **Index** value. This is their identification value! You'll notice that as soon as you press **Enter**, the values will be reflected in the **.xmi** tree.
- Now you need to set the **Next** and **Previous** EReferences which will make it possible to move cards through the box. Given that there are no partition before the first, set **Previous** values only for your second and third partition to that first partition. Similarly, only set the **Next** values for the first, and second partition to their respective 'next' partitions.
- In the same fashion, double click on each **Card** you created and modify their values. In particular, update the **Back** attribute to **One**. This is the value you'll be able to see from the partition and thus, the **.xmi** tree. You'll be experimenting with the **Face** attribute shortly, so provide a 1 value for that as well.
- Fill in the rest the cards with similar vocabulary-style words (such as two/2, three/3, ...) that you can 'test' yourself on (Figure 3.5). You now have a unique, customized learning box! Save your model, and ensure no errors exist before proceeding.

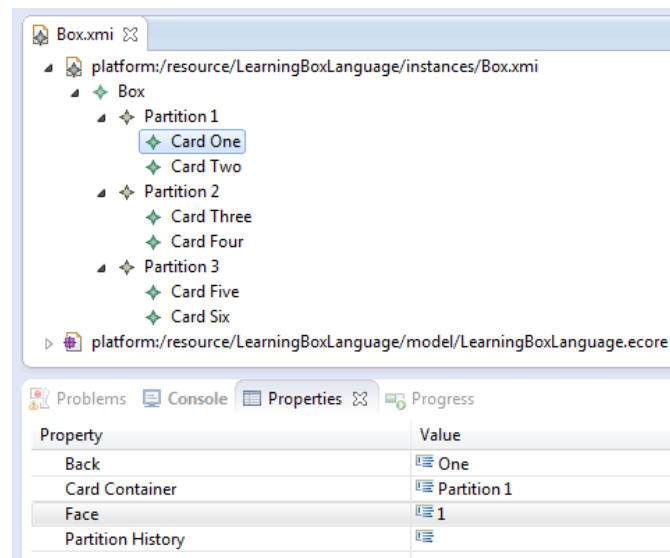


Figure 3.5: A vocabulary-style learning box

4 eMoflon’s graph viewer

At this point, you should now have a small instance model. While Eclipse’s built-in editor offers a nice tree structure and table to view and edit your model, wouldn’t it be nice to have a visualisation? eMoflon offers a “Graph View” to do exactly that! You’ll find that this is an effective feature to view how each element in your instance model interacts with others.

- ▶ When you first opened the eMoflon perspective, a small window to the right should have appeared with two tabs, “Outline” and “Graph View.”¹⁰
- ▶ Activate the second tab and drag one of your `Partition` instances into the empty window.
- ▶ To the right you can see the partition that was dragged in, double click on it to visualise all referenced objects from `Partition`. (Figure 4.1).
- ▶ Clicking any item in the view will highlight it, and hovering over a node will list all its properties in a pop-up dialogue. If you hover over

¹⁰If this window is not open, you can re-activate it by right-clicking on the “eMoflon” perspective in the main toolbar and pressing “Reset,” or by going to “Window/Show View/Other...” then “Other/Graph View”

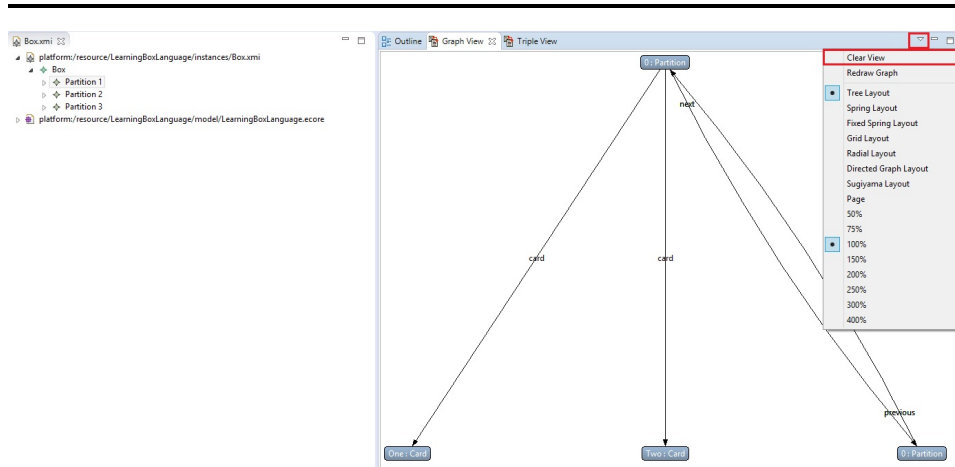


Figure 4.1: A single partition in eMoflon's graph view

a card, for example, you should be able to see their **back** and **face** values. You can also reposition elements.

- Try dragging a second **Partition** element from the model to the viewer. As you can see, it doesn't replace the current view, it simply places them side-by-side, showing the references between them partitions. To clear the screen, click the upside-down triangle in the top right corner of the window, and select "Clear View" (Figure 4.1).
- To make the graph bigger, increase the size of the window, then select "Redraw Graph" from the same upside-down arrow. As you can see, when an instance is already loaded, the graph view is not automatically updated. This option is useful if you change a property of an element (ie., the **back** value of a **Card**) and wish to see it reflected in the graph.
- Try experimenting with each of the different instance elements, viewer layouts and zoom settings found under the arrow. You'll notice for the "Spring Layout" that each time you press "Redraw Graph," the graph will re-arrange itself, even if you haven't updated any values or changed the window size.
- The "Graph View" features is not exclusively for instance models; Expand the second root node in the editor, and drag and drop the **LearningBoxLanguage** package into the viewer. Click on some or all nodes to open your entire metamodel completely. Now it is displayed in the viewer (Figure 4.2). We have found that the "Radial Layout" works best for a view of this complexity. This might be confusing at first, but the view contains your metamodel displayed in its abstract syntax, i.e., as an instance of Ecore. In contrast, the tree view displays the metamodel in UML-like concrete syntax.

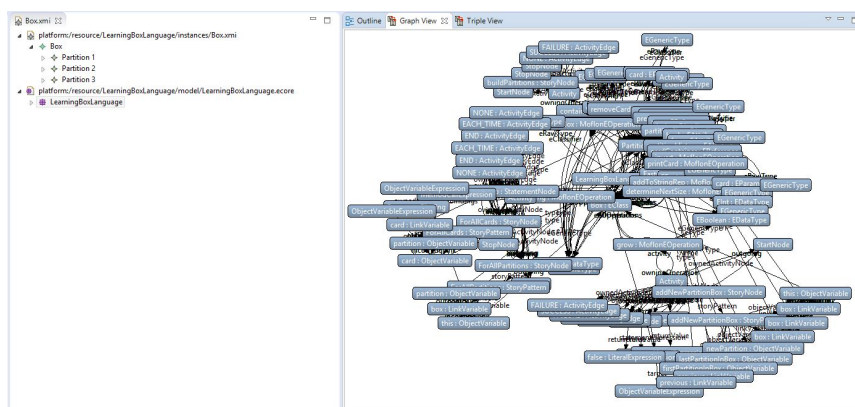


Figure 4.2: Our instance's complete type graph

5 Introduction to injections

This short introduction will show you how to implement small methods by adding handwritten code to classes generated from your model. Injections are inspired by partial classes in C#, and are our preferred way of providing a clean separation between generated and handwritten code.

Let's implement the `removeCard` method, declared in the `Partition` EClass. In order to 'remove' a card from a partition, all one needs to do is disable the link between them. Don't forget that (according to the signature) not only does `removeCard` have to pass in a `Card`, it must return one as well.

- From your working set, open "gen/LearningBoxLanguage.impl/PartitionImpl.java" and enter the following code in the `removeCard` declaration, starting at approximately line 347. Do not remove the first comment, which is necessary to indicate that this code is written by the user and needs to be extracted automatically as an injection. Please also do not copy and paste the following code – the copying process will most likely add invisible characters that eMoflon is unable to handle.

```
public Card removeCard(Card toBeRemovedCard) {
    // [user code injected with eMoflon]
    if(toBeRemovedCard != null){
        toBeRemovedCard.setCardContainer(null);
    }
    return toBeRemovedCard;
}
```

Figure 5.1: Implementation of `removeCard`

- Save the file, then right-click either on the file in the package explorer, or in the editor window, and choose "eMoflon/ Create/Update Injection for class" (Alt+Shift+E,I) from the context menu (Figure 5.2).
- This will create a new file in the "injection" folder of your project with the same package and name structure as the Java class, but with a new `.inject` extension (Figure 5.3).
- Double click to open and view this file. It contains the definition of a *partial class* (Figure 5.4).

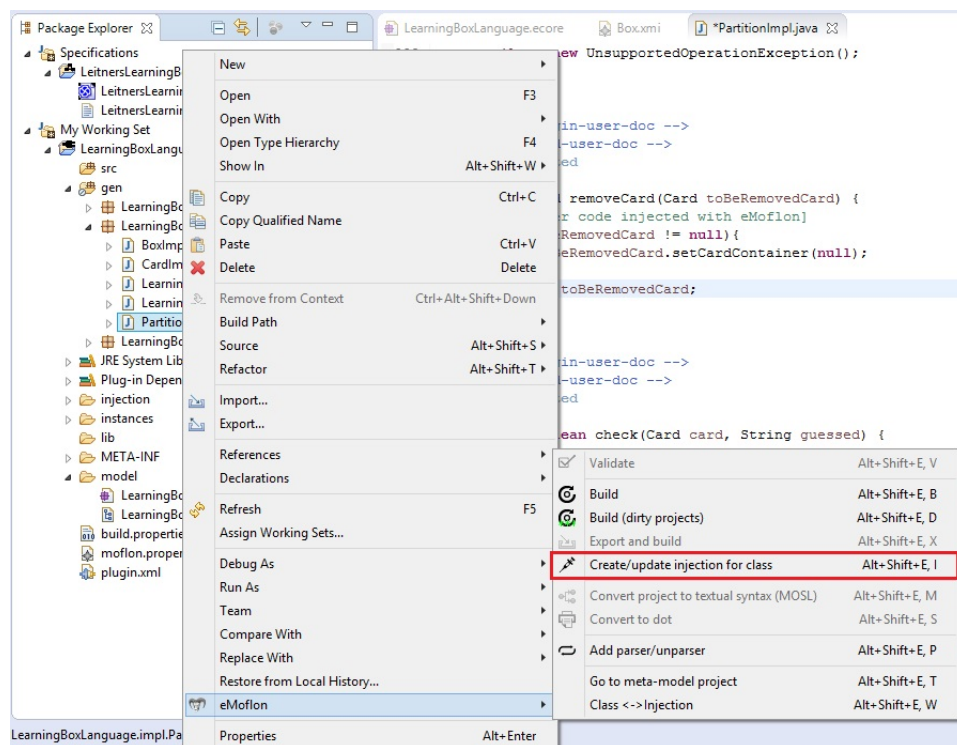


Figure 5.2: Create a new injection

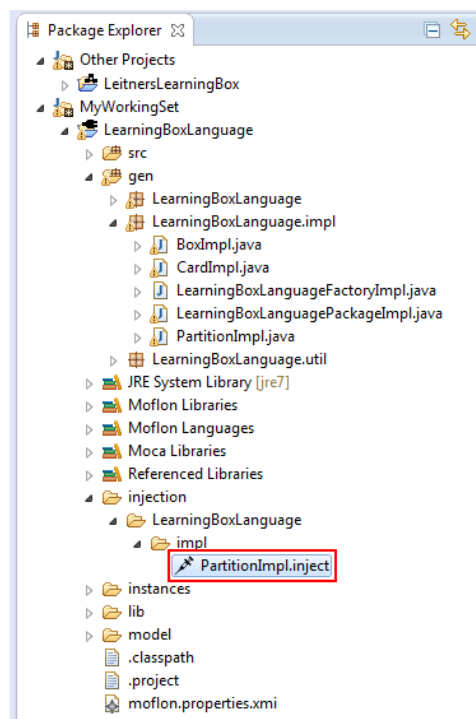


Figure 5.3: Partition injection file

The screenshot shows the content of the `PartitionImpl.inject` file. The code is as follows:

```
1 |
2 partial class PartitionImpl {
3
4
5
6 @model removeCard (Card toBeRemovedCard) <--
7
8     if(toBeRemovedCard != null){
9         toBeRemovedCard.setCardContainer (null);
10     }
11     return toBeRemovedCard;
12 -->
13
14 }
```

Figure 5.4: Generated injection file for PartitionImpl.java

-
- As a final step, build your metamodel to check that the code is generated and injected properly.
 - By the way, eMoflon allows you to switch quickly between a Java class and its injection file. When inside “PartitionImpl.java”, open the context menu and select “eMoflon/Class <-> Injection” (Alt+Shift+E,W). This brings you to “PartitionImpl.java”.

Repeat this command and you are back in “PartitionImpl.inject”.

- That’s it! While injecting handwritten code is a remarkably simple process, it is pretty boring and low level to call all those setters and getters yourself. We’ll return to injections for establishing two simple methods in Part III using this strategy, but we’ll also learn how to implement more complex methods using Story Diagrams.

Creating injections automatically with save actions

Information loss is a typical pitfall that you may encounter when working with injections: If you forget to create an injection and trigger a rebuild (“eMoflon/Build”, Alt+Shift+E,B), the generated code is entirely dropped – including any unsaved injection code.

To save you from this frustrating experience, eMoflon may automatically save injections whenever you save your Java file.

- Open the preferences dialog via “Window/Preferences” and navigate to the “Save Actions” (“Java/Editor/Save Actions”).
- To enable save actions, tick “Perform the selected actions on save” and “Additional actions” (Figure 5.5).
- Press “Configure”, switch to the “eMoflon Injections” tab and tick “Enable injection extraction on save” (Figure 5.6).
- After confirming the dialog, the bulleted list should contain the entry “Create eMoflon injections”.
- Open up “PartitionImpl.java”, perform a tiny modification on it, and save the file. The eMoflon console should display a message to confirm that injections have been extracted automatically, e.g.,

```
[handlers.CreateInjectionHandler::115] - Created injection
file for 'PartitionImpl'.
```

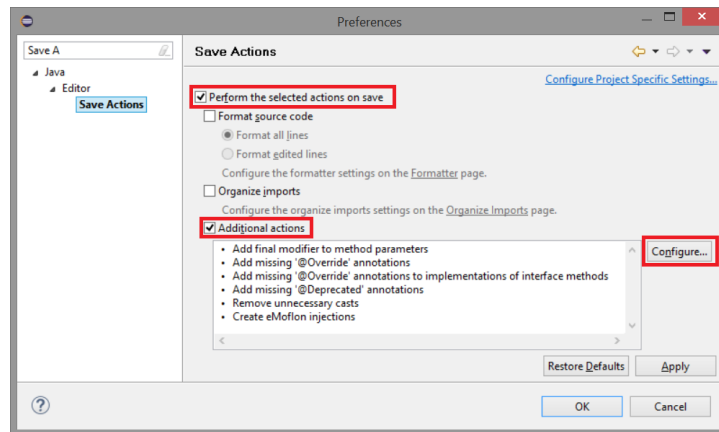


Figure 5.5: Main configuration dialog for save actions




Figure 5.6: Configuration dialog for additional save actions

6 Leitner's Box GUI

We would like to now provide you with a simple GUI with which you can take your model for a spin and see it in action.

- Navigate to the “Install, configure and deploy Moflon” button and open the “Install Workspace” menu bar (Figure 6.1).
- Load “Handbook Example (GUI)” (Figure 6.1). This will load the new project into your workspace. Right click `LeitnersBox` to invoke the context menu and select “Run as/Java application.”

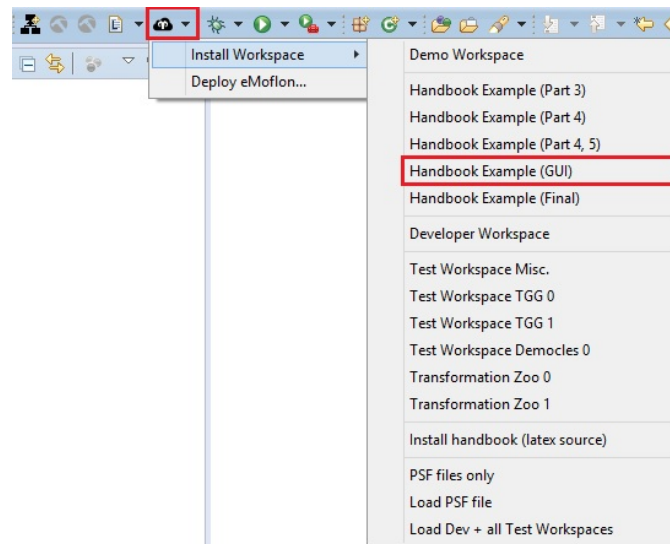


Figure 6.1: Load Leitner's Box GUI

- The GUI will automatically navigate to the instances folder where you've stored your instance model, then load your partitions and cards into the visualized box. Please note that this will only work if you named your dynamic instance `Box.xml` and placed it in the `instances` folder as suggested.
- Navigate to any card, and you'll be able to see two options, “Remove Card” and “Check Card” (Figure 6.2). While we'll implement “Check Card” in Part III, “Remove Card” is currently active, implemented by the injection you just created.

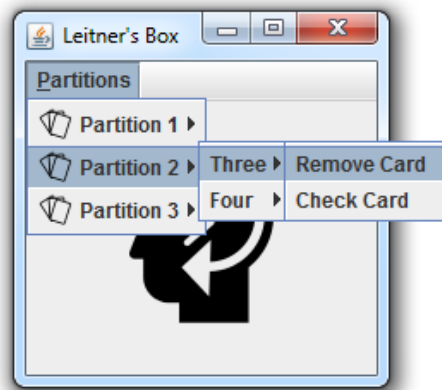


Figure 6.2: Using the GUI with your instance

- Experiment with your instance model and confirm your injection works by removing some items from the GUI. You'll notice the change immediately in the `Box.xml` file. You can also close the GUI and add, remove, or rename more elements in the model, then observe how the changes are reflected in the GUI.
- Expand the "Other Projects" node and explore `LeitnersBoxController` and `LeitnersBoxView` files to get an understanding of the GUI. Can you find where the controller loads and connects to the model?

7 Conclusion and next steps

Whoo, this has been quite a busy handbook. Great job, you've finished Part II! This part contained some key skills for eMoflon, as we learned how to create the abstract syntax in Ecore for our learning box! The specification is now complete with all the classes, attributes, references, and method signatures that make up the type graph for a working learning box. Additionally, we also learned how to insert a small handwritten method into generated code, and tested all our work in an interactive GUI.

If you enjoyed this section and wish to fully develop *all* the methods we just declared, we invite you to carry on to Part III: Story Diagrams¹¹! Story Diagrams are a powerful feature of eMoflon as we can model a large part of a system's dynamic semantics via high-level pattern rules.

Of course, you're always free to pick a different part of the handbook if you feel like skipping ahead and checking out Triple Graph Grammars (TGGs) in Part IV¹². We'll provide instructions on how to easily download all the required resources so you can start without having to complete the previous parts.

For a detailed description of all parts, please refer to Part 0¹³.

Cheers!

¹¹Download: <https://emoflon.github.io/eclipse-plugin/beta/handbook/part3.pdf>

¹²Download: <https://emoflon.github.io/eclipse-plugin/beta/handbook/part4.pdf>

¹³Download: <https://emoflon.github.io/eclipse-plugin/beta/handbook/part0.pdf>

Glossary

Abstract Syntax Defines the valid static structure of members of a language.

Concrete Syntax How members of a language are represented. This is often done textually or visually.

Constraint Language Typically used to specify complex constraints (as part of the static semantics of a language) that cannot be expressed in a metamodel.

Dynamic Semantics Defines the dynamic behaviour for members of a language.

Grammar A set of rules that can be used to generate a language.

Graph Grammar A grammar that describes a graph language. This can be used instead of a metamodel or type graph to define the abstract syntax of a language.

Meta-Language A language that can be used to define another language.

Meta-metamodel A *modeling language* for specifying metamodels.

Metamodel Defines the abstract syntax of a language including some aspects of the static semantics such as multiplicities.

Model Graphs which conform to some metamodel.

Modelling Language Used to specify languages. Typically contains concepts such as classes and connections between classes.

Static Semantics Constraints members of a language must obey in addition to being conform to the abstract syntax of the language.

Type Graph The graph that defines all types and relations that form a language. Equivalent to a metamodel but without any static semantics.

Unification An extension of the object oriented “Everything is an object” principle, where everything is regarded as a model, even the metamodel which defines other models.